

Reconfigurable Computing

Tutorial 1 - LSTM Network

Philip Leong

September 7, 2017

1 Introduction

The long short-term memory (LSTM) network has been used with great success in time-series prediction problems. In this series of tutorials, we will develop an FPGA implementation of an LSTM. This tutorial has the following goals:

- Practise translating published mathematical descriptions into hardware designs
- Gain experience in using high level synthesis
- Gain experience in design optimisation.

Our implementation will have identical functionality to BasicLSTMCell¹ which is used in Google's Tensorflow neural network software. This, in turn, is an implementation of a regularisation scheme published in reference [1].

Using the same notation as [1], let $h_t^l \in \mathbb{R}^{n_l}$ be a hidden state in layer l at timestep t , $T_{m,n} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ be an affine transformation ($Wx + b$ for some $W \in \mathbb{R}$ and b), \odot elementwise multiplication, sigm is the elementwise sigmoid function, \tanh is the elementwise hyperbolic tangent, and h_t^0 be an input vector at timestep t . We use the activations h_t^L to predict y_t where L is the number of layers in the LSTM.

The neural network implements a state transition

$$LSTM : h_t^{l-1}, h_{t-1}^l, c_{t-1}^l \rightarrow h_t^l, c_t^l \quad (1)$$

where

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{sigm} \\ \tanh \end{pmatrix} = T_{(n_{l-1}+n_l), (4n_l)} \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix} \quad (2)$$

$$c_t^l = f \odot c_{t-1}^l + i \odot g \quad (3)$$

$$h_t^l = o \odot \tanh(c_t^l) \quad (4)$$

¹<https://github.com/tensorflow/tensorflow/blob/master/tensorflow/python/ops/rnn.py>

2 Laboratory Questions

In answering these questions, marks will be awarded not only for correctness but also understandability and elegance of the solution.

2.1 (20%) Familiarisation with `xsimple.py`

Download the tutorial source files using the command

```
git clone git@github.com:phwl/hlslstm.git
```

You can check that everything is working by typing:

```
cd hlslstm/src
make test
```

You should see output similar to the below.

```
python xsimple.py >xsimple.out
g++ -p -o simple simple.cpp -lm
grep pred xsimple.out > /tmp/xsimple.out
echo "testing simple ..."
testing simple ...
(./simple | grep pred | sdiff -w80 /tmp/xsimple.out -)
vprint(y_pred) -0.0055398695 -0.01926    vprint(y_pred) -0.0055398695 -0.01926
vprint(y_pred) -0.012691636 -0.027500    vprint(y_pred) -0.012691636 -0.027500
vprint(y_pred) -0.019167556 -0.028273    vprint(y_pred) -0.019167556 -0.028273
vprint(y_pred) -0.021765375 -0.026982 | vprint(y_pred) -0.021765375 -0.026983
rm -f /tmp/xsimple.out
```

The output is the result of an `sdiff` command which compares the `y_pred` output of a floating-point python implementation of the LSTM (on the left hand side) with the output of the `gen.cpp` C program (on the right hand side). If the lines of text are different (as in the case in the last line), a `|` appears between them. Note that the simple program prints a number of variables, but (make `test`) only compares the `y_pred` outputs.

Change the program so it accepts vectors of length 32 and produces output predictions of length 32.

2.2 (40%) Testbench

Also modify the program so it has a main program (which we will call the testbench), that checks if the relative error of each individual output is less than 10%. The program should print "PASS" and the mean squared error (MSE) if correct. Otherwise it should print a diagnostic test message indicating which test pattern and output is in error (the expected output is provided in the `l_y[]` array defined in `gen\io.h`). Change the data type used for calculation from double to float. Rerun and see if it passes the testbench. Is the MSE the same?

2.3 (40%) Vivado HLS Project

Create a Vivado HLS project² and import the `simple.cpp` program. Generate a synthesis report for the project. How many cycles does the execution take? What is the clock period?

References

- [1] Wojciech Zaremba, Ilya Sutskever, and Oriol Vinyals. Recurrent neural network regularization. *CoRR*, abs/1409.2329, 2014.

²https://www.xilinx.com/support/documentation/sw_manuals/xilinx2015_4/ug871-vivado-high-level-synthesis-tutorial.pdf